

# MAC - the medical appointment coordinator (A chatbot prototype)

Karen Gaffney 1.20.26



# Why use a chatbot for medical appointment scheduling?

- **Single interface** to schedule appointments, no need to navigate a menu
- **Relative intent** Chatbots understand “next tuesday”
- **Data and Insights** within the chat “when was the last time Mary came in for an exam?”
- **Accessibility** allows voice to text integration can offer hands free scheduling for busy doctors




# Hi I'm Mac! (Medical Appointment Coordinator)

A Chatbot for **Doctors** to schedule appointments with their Patients

- Schedule appointments
  - Capture intent, gather required information, and then ask for confirmation before scheduling the appointment
  - Patient email is retrieved from SQL database by name
  - Redirect to add\_patient if patient is not in database
- Adds new patients
  - Capture intent, gather required information, and then ask for confirmation before adding a patient to the patient database

# Tech Stack






# LLM workflow Trial and Error 1

1. Prompt AI with Doctor's message and system prompt to **return structured json** containing appointment details
  - a. Doesn't see previous messages, all information must be in one prompt

```
{  
    "assistant_message": string,  
    "draft_event": {  
        "attendee_name": string | null,  
        "start_time_local": string | null,  
        "duration_minutes": number | null,  
        "title": string | null,  
        "notes": string | null  
    }  
}
```






# LLM workflow Trial and Error 2

Prompt AI with a history of Doctor messages, AI response, and system prompt to **return structured json** containing appointment details.

Doesn't ask for already provided info

- If all required info is present, overwrite ai response with predefined confirmation message, execute event when `confirm_intent = True`
- Breaks json structure after 3~4 messages (style drift)






# Final LLM-assisted workflow


1. **State based updating.** Front end generates a **unique session ID** and sends it with every Doctor's message. Session ID retrieves an existing state or initializes a new state.
2. AI receives the **current state + latest Doctor's message**
  - a. **State** tracks
    - i. Current mode [schedule/add]
    - ii. Awaiting user confirmation
    - iii. Draft event
  - b. **AI returns**
    - i. Assistant message
    - ii. Updates to draft event
    - iii. Changes in intent from current state
    - iv. Confirmation intent

```
{"draft_event": {  
  "attendee_name": None,  
  ...  
  "notes": None},  
  "awaiting_confirmation": False,  
  "mode": "schedule",  
  "pending_patient": {  
    "name": None  
    ...  
    "dob": None  
  }  
}  
  
{  
  "assistant_message": string,  
  "updates": {  
    "attendee_name": ...  
    "intent": null | schedule | add_patient  
  },  
  "confirm_intent": "yes" | "no" | "unknown"  
}
```

# Demo Video



# System Design





# Strengths of Design

- Avoids hallucinated actions - events must be awaiting confirmation and be confirmed
  - **Awaiting\_confirmation** = set by back end logic when all required fields are filled
  - **Confirm\_intent** = intent parsed by LLM on user response to confirm
- State object and updates workflow allow for clearer debugging
- Makes conversations resumable and able to switch between intents
- Scales cleanly to new intents and workflows
  - Add new intent options beyond add\_patient and schedule



# Testing and Error Handling

- Test routes for Microsoft Graph API functions
  - Manually entered payloads and events, verified execution in Microsoft Calendar
- Logging of state and LLM response for each unique session
  - Draft event
  - Awaiting confirmation status
  - LLM raw response
    - Update json
    - Confirm\_intent
    - Assistant message
- Future testing for handling non json responses, connection timeouts to SQL server



# Future Extensibility

- Calendar conflict handling
  - Give the chatbot a view of the Doctor's calendar in a X day window to offer alternative times should there be a conflicting meeting
- Login for multiple Doctors (HIPPA privacy considerations)
  - Relational Database Schema would allow for Doctors to schedule appointments with their patients only, and retrieve information about their patients only
  - Login would use microsoft entra authentication to meet security standards
- Salting and hashing sensitive patient data - DOB, SSN etc
- Front End Features
  - Database view on past appointments (notes)
  - Calendar view on scheduled appointments
  - Database view on patient data with filtering based on Logged In Doctor

# Prototype vs Production

- LLM
  - Using GPT 5 mini for speed and latency
  - Evaluate multiple models for different tasks. Add RAG augmented models for domain specific tasks. Semantic Kernel for schema stability
- Secrets
  - .env, loaded in to container at runtime
  - Migrate secrets to Azure Vault
- Identity/Tenant management
  - Single tenant support for Azure Graph
  - Multi-tenant architecture with client Entra ID
- State management
  - In memory session by front end random UUID. Past states retained indefinitely for debugging (till container is reloaded) <-memory leak
  - Move session state to Azure Cache. Implement TTL-based expiration
- Data Layer
  - Azure SQL + mssql for query execution
  - Azure SQL + mssql for transactional work. Databricks/Fabric for analytics, reporting, cross-tenant insights
- Deployment
  - Locally hosted with manual firewall IP access to SQL server
  - Deploy to Azure with CI/CD pipelines and comprehensive tests. No public IP firewall rules, Azure integrated networking



# Collaboration with others

- User specifications and feedback
  - Doctor input on prioritization of features
  - Testing of real world performance
- Division of work
  - Front End UI
    - Pages for calendar view, patient data view
  - Back End
    - New intents + actions. System prompts, external API integration, build event payloads
  - Database
    - Relational Schemas for patient/doctor. New tables to store past patient notes
- IT Security and HIPPA standards
  - Microsoft Entra ID integration for more doctors
  - Sufficient encryption of patient data



# Thank you! Any Questions?

MAC is a example prototype that shows properly structured AI workflows can translate **human intent** into **real-world actions** without sacrificing security or transparency